
Confluence Documentation

AWS Architecture

November 16, 2021

Nikki Tebaldi (ntebaldi@umass.edu)

TABLE OF CONTENTS

| | |
|---------------------------------|---|
| Table Of Figures | 2 |
| AWS Architecture Overview | 3 |
| VPC..... | 3 |
| IAM Resources | 4 |
| Compute Environments..... | 4 |
| Job Queues..... | 4 |
| Job Definitions | 5 |
| ECR Containers..... | 6 |
| EFS..... | 7 |
| S3 Buckets..... | 7 |
| List of Acronyms..... | 8 |

TABLE OF FIGURES

| | |
|--------------------------------|---|
| Figure 1 AWS Architecture..... | 3 |
|--------------------------------|---|

AWS ARCHITECTURE OVERVIEW

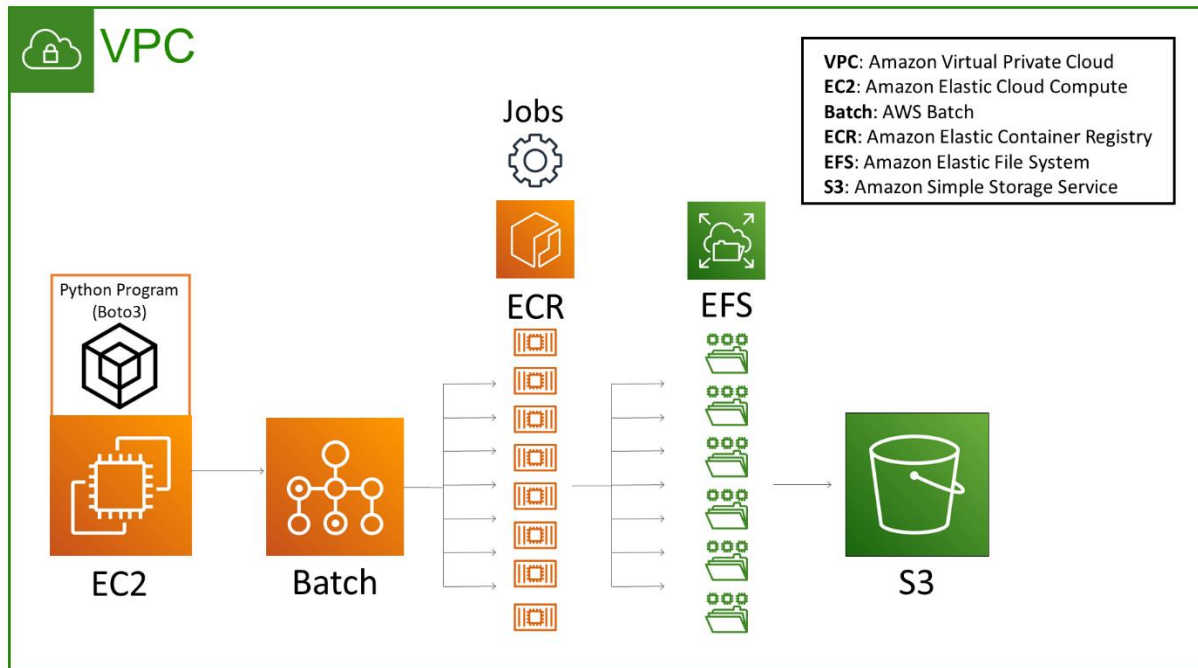


Figure 1 AWS Architecture

Confluence as a platform operates within several AWS architectural components. AWS Batch is used to submit and execute containerized jobs for each module in a stage. The containers are hosted in the AWS Elastic Container Registry (ECR). The Batch infrastructure is made up of several compute environments that are backed by AWS Fargate. A job queue exists for each stage and job definitions are created for each module. Intermediate storage is provided through Amazon Elastic File System (EFS) and is provided for each stage of processing. The results of an execution of the Confluence workflow and prior data information is stored in an Amazon Simple Storage Service (S3) bucket.

VPC

| Name | Description | Region | Subnets in Availability Zones |
|------------|--|--------------------|---|
| confluence | Virtual Private Cloud network to execute Batch jobs in | us-west-2 (Oregon) | us-west-2a (private), us-west-2b (private), us-west-2c (private and public) |

IAM RESOURCES

| Name | Description | Component Notes |
|-----------------------------|---|---|
| ecsTaskExecutionRole | Batch Fargate Role that allows ECS task execution | Grants ECS container and Fargate agents to make AWS API calls |
| BatchEFSJobRole | Batch Fargate Role that allows job access to EFS | Allows ECS container access to EFS |

COMPUTE ENVIRONMENTS

| Name | Description | Max vCPUs |
|--------------------|--------------------------------------|-----------|
| data | Managed, Fargate compute environment | 1000 |
| diagnostics | Managed, Fargate compute environment | 1000 |
| flpe | Managed, Fargate compute environment | 1000 |
| moi | Managed, Fargate compute environment | 1000 |
| offline | Managed, Fargate compute environment | 1000 |
| validation | Managed, Fargate compute environment | 1000 |

JOB QUEUES

| Name | Compute environment |
|-----------------------------|---------------------|
| input | data |
| flpe | flpe |
| moi | moi |
| priors | data |
| prediagnostics | diagnostics |
| postdiagnostics-flpe | diagnostics |
| postdiagnostics-moi | diagnostics |
| offline | offline |
| output | data |
| validation | validation |

JOB DEFINITIONS

| Name | Memory | vCPU | EFS Volume: Mount Point |
|---------------------|--------|------|---|
| geobam | 8192 | 4 | input: /mnt/data/input flpe: /mnt/data/output |
| hivdi | 4096 | 2 | input: /mnt/data/input flpe: /mnt/data/output |
| metroman | 8192 | 4 | input: /mnt/data/input flpe: /mnt/data/output |
| momma | 0.25 | 512 | input: /mnt/data/input flpe: /mnt/data/output |
| sad | 2048 | 1 | input: /mnt/data/input flpe: /mnt/data/output |
| sic4dvar | 2048 | 1 | input: /mnt/data/input flpe: /mnt/data/output |
| input | 8192 | 1 | input: /mnt/data |
| moi | 2048 | 1 | input: /mnt/data/input flpe: /mnt/data/flpe moi: /mnt/data/output |
| offline | 2048 | 1 | input: /mnt/data/input offline: /mnt/data/output |
| output | 2048 | 1 | input: /mnt/data/input flpe: /mnt/data/output moi: /mnt/data/moi diagnostics: /mnt/data/diagnostics offline: /mnt/data/offline validation: /mnt/data/validation output: /mnt/data/output |
| postdiagnostic_flpe | 1024 | 0.5 | input: /mnt/data/input flpe: /mnt/data/flpe diagnostics/postdiagnostics/reach: /mnt/data/output |
| postdiagnostics_moi | 1024 | 0.5 | input: /mnt/data/input flpe: /mnt/data/flpe moi: /mnt/data/moi |

| | | | |
|-----------------------|------|-----|---|
| | | | diagnostics/postdiagnostics/basin: /mnt/data/output |
| prediagnostics | 1024 | 0.5 | input: /mnt/data/input diagnostics/prediagnostics: /mnt/data/output |
| priors | 8192 | 1 | input: /mnt/data |
| validation | 2048 | 1 | input: /mnt/data/input offline: /mnt/data/offline validation: /mnt/data/output |

ECR CONTAINERS

| Name | URI | Visibility | Dockerfile Link | Size (MB) |
|-----------------------|---|------------|--|-----------|
| hivdi | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/hivdi | private | HiVDI | 1542.66 |
| input | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/input | private | Input | 339.07 |
| metroman | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/metroman | private | MetroMan | 185.91 |
| moi | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/moi | private | MOI | 403.75 |
| momma | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/momma | private | MOMMA | 440.61 |
| offline | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/offline | private | Offline | 93.27 |
| output | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/output | private | Output | 115.93 |
| postd-flpe | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/postd-flpe | private | Postdiagnostics FLPE | 403.75 |
| postd-moi | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/postd-moi | private | Postdiagnostics MOI | 402.71 |
| prediagnostics | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/prediagnostics | private | Prediagnostics | 344.59 |
| priors | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/priors | private | Priors | 1170.18 |
| sic4dvar | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/sic4dvar | private | n/a | 273.96 |
| validation | xxxxxxxxxxxx.dkr.ecr.us-west-2.amazonaws.com/validation | private | Validation | 176.94 |
| geobam | public.ecr.aws/s8t4s7g5/geobam | public | geoBAM | 590.51 |
| sad | public.ecr.aws/s8t4s7g5/sad | public | SAD | 314.47 |

EFS

| Name | ID | Contents |
|-------------|-------------|---|
| input | fs-60e5fd67 | basin.json continent.json reaches.json sets.json gage/ sos/ sword/ swot/ |
| flpe | fs-b6f6eeb1 | hivdi/ geobam/ metroman/ momma/ sad/ sic4dvar/ |
| moi | fs-89c6048f | |
| diagnostics | fs-92ea2894 | prediagnostics/ postdiagnostics/basin/ postdiagnostics/reach/ |
| offline | fs-39f13f3f | |
| validation | fs-b55580b3 | figs/ stats/ |
| output | fs-ea26ebec | version/ |
| logs | fs-1d71a91b | sic4dvar/ |

S3 BUCKETS

| Name | URI | Region | Contents |
|-----------------|----------------------|----------------------------|---|
| confluence-sos | s3://confluence-sos | US West (Oregon) us-west-2 | constrained/ figs/ unconstrained/ |
| confluence-swot | s3://confluence-swot | US West (Oregon) us-west-2 | pass249/ pass264/ pass527/ |

LIST OF ACRONYMS

AWS: Amazon Web Services

EC2: Amazon Elastic Compute Cloud

ECR: Amazon Elastic Container Registry

EFS: Amazon Elastic File System

FLPE: Flow Law Parameter Estimation

IAM: AWS Identity and Access Management (IAM)

JPL: Jet Propulsion Laboratory

PO.DAAC: Physical Oceanography Distributed Active Archive Center

S3: Amazon Simple Storage Service

SDS: Science Data System

SoS: SWORD of Science

SWORD: SWOT River Database

SWOT: Surface Water and Ocean Topography

VPC: Amazon Virtual Private Cloud